

Process of enabling kubectl logs feature

1. Make sure the **ca.crt** and **ca.key** files are in the kubernetes directory. If the kubernetes cluster is built up by kubeadm. Then, check the directory with the following command:
`ls /etc/kubernetes/pki/`
2. Set CLOUDCOREIPS env (must be created as root). This environment variable should be the cloudcore IP address.
 Setting CLOUDCOREIPS needs the following command:
`export CLOUDCOREIPS="192.168.0.139"`
 (Warning: we must use the same "terminal" to continue our work, or we need to type this command again.)
 Checking the environment variable with the following command:
`echo $CLOUDCOREIPS`
3. Generate the certificates for CloudStream on cloud node, however, the generation file is not in the /etc/kubeedge/, we need to copy it from the repository which was git cloned from GitHub.
 Change user to root:
`sudo su`
 Copy certificates generation file from original cloned repository:
`cp $GOPATH/src/github.com/kubeedge/kubeedge/build/tools/certgen.sh /etc/kubeedge/`
 Change directory to the kubeedge directory:
`cd /etc/kubeedge/`
 Generate certificates from **certgen.sh**
`certgen.sh stream`
4. It is needed to set iptables or firewall on the host. (It is referred to the KubeEdge document.)
 Setting up the iptables for CloudStream (Port 10003 and 10350 are the default ports for the CloudStream and edgecore.):
 (This command should be executed on every apiserver deployed node.)
 (In my case, this is my master node, and execute this command by root.)
`iptables -t nat -A OUTPUT -p tcp --dport 10350 -j DNAT --to $CLOUDCOREIPS:10003`
 Also, we need to clean up the iptables on Raspberry Pi 4 (Edge node) to make **kubectll logs** work.
`iptables -F && iptables -t nat -F && iptables -t mangle -F && iptables -X`

Process of enabling kubectl logs feature

1. Make sure the **ca.crt** and **ca.key** files are in the kubernetes directory. If the kubernetes cluster is built up by kubeadm. Then, check the directory with the following command:

```
ls /etc/kubernetes/pki/
```

2. Set CLOUDCOREIPS env (must be created as root). This environment variable should be the cloudcore IP address.

Setting CLOUDCOREIPS needs the following command:

```
export CLOUDCOREIPS="192.168.0.139"
```

(Warning: we must use the same "terminal" to continue our work, or we need to type this command again.)

Checking the environment variable with the following command:

```
echo $CLOUDCOREIPS
```

3. Generate the certificates for CloudStream on cloud node, however, the generation file is not in the /etc/kubeedge/, we need to copy it from the repository which was git cloned from GitHub.

Change user to root:

```
sudo su
```

Copy certificates generation file from original cloned repository:

```
cp $GOPATH/src/github.com/kubeedge/kubeedge/build/tools/certgen.sh /etc/kubeedge/
```

Change directory to the kubeedge directory:

```
cd /etc/kubeedge/
```

Generate certificates from **certgen.sh**

```
certgen.sh stream
```

4. It is needed to set iptables or firewall on the host. (It is referred to the KubeEdge document.)

Setting up the iptables for CloudStream (Port 10003 and 10350 are the default ports for the CloudStream and edgecore.):

(This command should be executed on every apiserver deployed node.)

(In my case, this is my master node, and execute this command by root.)

```
iptables -t nat -A OUTPUT -p tcp --dport 10350 -j DNAT --to $CLOUDCOREIPS:10003
```

Also, we need to clean up the iptables on Raspberry Pi 4 (Edge node) to make kubectl logs work.

```
iptables -F && iptables -t nat -F && iptables -t mangle -F && iptables -X
```

5. Modify "both" **/etc/kubeedge/config/cloudcore.yaml** and **/etc/kubeedge/config/edgecore.yaml** on cloudcore and edgecore. Set up cloudStream and edgeStream to enable: true.

Also, you have to change the server IP to your cloudcore IP (the same as \$CLOUDCOREIPS).

Open the YAML file in cloudcore:

```
sudo nano /etc/kubeedge/config/cloudcore.yaml
```

Modify the file in the following part:

```
cloudStream:
  enable: true
  streamPort: 10003
  tlsStreamCAFile: /etc/kubeedge/ca/streamCA.crt
  tlsStreamCertFile: /etc/kubeedge/certs/stream.crt
  tlsStreamPrivateKeyFile: /etc/kubeedge/certs/stream.key
  tlsTunnelCAFile: /etc/kubeedge/ca/rootCA.crt
  tlsTunnelCertFile: /etc/kubeedge/certs/server.crt
  tlsTunnelPrivateKeyFile: /etc/kubeedge/certs/server.key
  tunnelPort: 10004
```

Open the YAML file in edgecore:

```
sudo nano /etc/kubeedge/config/edgecore.yaml
```

Modify the file in the following part:

```
edgeStream:
  enable: true
  handshakeTimeout: 30
  readDeadline: 15
  server: 192.168.0.139:10004
  tlsTunnelCAFile: /etc/kubeedge/ca/rootCA.crt
  tlsTunnelCertFile: /etc/kubeedge/certs/server.crt
  tlsTunnelPrivateKeyFile: /etc/kubeedge/certs/server.key
  writeDeadline: 15
```

6. Restart all the cloudcore and edgecore

*Warning: there are some containers which are related to the edgecore. Thus, we have to restart edgecore.service first and make those kube-proxy related containers “pending” or “pause”, and then kill those containers. Finally, we can restart edgecore again to make it work properly.

Before you start to restart any service, please turn your user into root. Since you have to possess the rights of accessing the docker and system services directly.

`sudo su`

cloudCore:

`pkill cloudcore`

`nohup cloudcore > cloudcore.log 2>&1 &`

edgeCore:

`systemctl restart edgecore.service`

`docker ps`

```
jroot@raspberrypi:/home/pi# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	
AMES						N
d40ad407d6f9	kubeedge/pause-arm:3.1	"/pause"	12 minutes ago	Up 12 minutes		k
8s_POD_weave-net-fr9vx_kube-system_b30flab0-8cf2-448a-8588-ad191f5f9b37_13						
f59bdacbf48	b732e226e0f1	"/usr/local/bin/kube.."	12 minutes ago	Up 12 minutes		k
8s_kube-proxy_kube-proxy-s8476_kube-system_5c0e5bc4-73ae-4daa-a0fd-d8d8713ddad5_14						
f15a4b9aae7a	kubeedge/pause-arm:3.1	"/pause"	13 minutes ago	Up 13 minutes		k
8s_POD_kube-proxy-s8476_kube-system_5c0e5bc4-73ae-4daa-a0fd-d8d8713ddad5_13						
2559412c2401	c1fd304a4885	"temperature-mapper"	13 minutes ago	Up 13 minutes		k
8s_temperature_temperature-mapper-79f4989f65-5tszw_default_867db8c5-f606-45ca-9f61-7912fe4df84d_0						
38ac0fc561b7	kubeedge/pause-arm:3.1	"/pause"	13 minutes ago	Up 13 minutes		k
8s_POD_temperature-mapper-79f4989f65-5tszw_default_867db8c5-f606-45ca-9f61-7912fe4df84d_0						

`docker container kill d40ad407d6f9`

`docker container kill f59bdacbf48`

`docker container kill f15a4b9aae7a`

`systemctl restart edgecore.service`

Use the following command to check if **kubectrl logs** feature is activated or not on localhost (The Raspberry Pi 4 is the KubeEdge Edge node in my case).

Check if the edge stream feature is activated or not:

`sudo netstat -tulpn | grep edgecore`

```
pi@raspberrypi:~$ sudo netstat -tulpn | grep edgecore
```

tcp	0	0	172.17.0.1:4000	0.0.0.0:*	LISTEN	2613/edgecore
tcp	0	0	127.0.0.1:10350	0.0.0.0:*	LISTEN	2613/edgecore
udp	0	0	172.17.0.1:53	0.0.0.0:*		2613/edgecore

Make sure if the your edge stream service can be accessed:

`curl "http://localhost:10350/stats/summary?only_cpu_and_memory=true"`

If you successfully activate the feature, you will see the following information:

```

root@raspberrypi:/home/pi# curl "http://localhost:10350/stats/summary?only_cpu_and_memory=true"
{
  "node": {
    "nodeName": "raspberrypi",
    "systemContainers": [
      {
        "name": "pods",
        "startTime": "2020-08-22T11:04:10Z",
        "cpu": {
          "time": "2020-08-23T02:29:35Z",
          "usageNanoCores": 12480180,
          "usageCoreNanoSeconds": 846749998105
        },
        "memory": {
          "time": "2020-08-23T02:29:35Z",
          "availableBytes": 8170373120,
          "usageBytes": 111263744,
          "workingSetBytes": 88150016,
          "rssBytes": 11223040,
          "pageFaults": 0,
          "majorPageFaults": 0
        }
      },
      {
        "name": "kubelet",
        "startTime": "2020-08-22T11:06:28Z",
        "cpu": {

```

```

    "time": "2020-08-23T02:29:34Z",
    "usageNanoCores": 72278481,
    "usageCoreNanoSeconds": 4337730875230
  },
  "memory": {
    "time": "2020-08-23T02:29:34Z",
    "availableBytes": 8796031279104,
    "usageBytes": 63361024,
    "workingSetBytes": 61739008,
    "rssBytes": 58212352,
    "pageFaults": 25766202,
    "majorPageFaults": 0
  }
},
{
  "name": "runtime",
  "startTime": "2020-08-22T11:04:09Z",
  "cpu": {
    "time": "2020-08-23T02:29:34Z",
    "usageNanoCores": 36197490,
    "usageCoreNanoSeconds": 1656552133365
  },
  "memory": {
    "time": "2020-08-23T02:29:34Z",
    "availableBytes": 8796035026944,
    "usageBytes": 120844288,
    "workingSetBytes": 57991168,

```

```

    "rssBytes": 39063552,
    "pageFaults": 12661506,
    "majorPageFaults": 198
  }
}
],
"startTime": "2020-08-22T11:04:09Z",
"cpu": {
  "time": "2020-08-23T02:29:40Z",
  "usageNanoCores": 166651877,
  "usageCoreNanoSeconds": 10859131447776
},
"memory": {
  "time": "2020-08-23T02:29:40Z",
  "availableBytes": 7455789056,
  "usageBytes": 1336512512,
  "workingSetBytes": 802734080,
  "rssBytes": 150327296,
  "pageFaults": 39732,
  "majorPageFaults": 66
}
},
"pods": [
  {
    "podRef": {
      "name": "kube-proxy-s8476",
      "namespace": "kube-system",

```

```

"uid": "5c0e5bc4-73ae-4daa-a0fd-d8d8713ddad5"
},
"startTime": "2020-08-22T11:53:32Z",
"containers": [
{
  "name": "kube-proxy",
  "startTime": "2020-08-22T11:53:43Z",
  "cpu": {
    "time": "2020-08-23T02:29:43Z",
    "usageNanoCores": 1063436,
    "usageCoreNanoSeconds": 103051749819
  },
  "memory": {
    "time": "2020-08-23T02:29:43Z",
    "availableBytes": 8796085219328,
    "usageBytes": 9015296,
    "workingSetBytes": 7798784,
    "rssBytes": 6275072,
    "pageFaults": 754347,
    "majorPageFaults": 0
  }
}
],
"cpu": {
  "time": "2020-08-23T02:29:42Z",
  "usageNanoCores": 1099277,
  "usageCoreNanoSeconds": 111640636254
}

```



```

},
"memory": {
  "time": "2020-08-23T02:29:42Z",
  "availableBytes": 8796056662016,
  "usageBytes": 42168320,
  "workingSetBytes": 36356096,
  "rssBytes": 6246400,
  "pageFaults": 0,
  "majorPageFaults": 0
}
},
{
  "podRef": {
    "name": "weave-net-fr9vx",
    "namespace": "kube-system",
    "uid": "b30f1ab0-8cf2-448a-8588-ad191f5f9b37"
  },
  "startTime": "2020-08-22T11:53:59Z",
  "containers": null,
  "cpu": {
    "time": "2020-08-23T02:29:32Z",
    "usageNanoCores": 0,
    "usageCoreNanoSeconds": 45122404150
  },
  "memory": {
    "time": "2020-08-23T02:29:32Z",
    "availableBytes": 8796046766080,

```

```

"usageBytes": 57335808,
"workingSetBytes": 46252032,
"rssBytes": 0,
"pageFaults": 0,
"majorPageFaults": 0
}
},
{
  "podRef": {
    "name": "temperature-mapper-79f4989f65-5tswz",
    "namespace": "default",
    "uid": "867db8c5-f606-45ca-9f61-7912fe4df84d"
  },
  "startTime": "2020-08-22T11:53:10Z",
  "containers": [
    {
      "name": "temperature",
      "startTime": "2020-08-22T11:53:12Z",
      "cpu": {
        "time": "2020-08-23T02:29:42Z",
        "usageNanoCores": 11894874,
        "usageCoreNanoSeconds": 654512498707
      },
      "memory": {
        "time": "2020-08-23T02:29:42Z",
        "availableBytes": 8796088307712,
        "usageBytes": 7143424,

```

```

    "workingSetBytes": 4710400,
    "rssBytes": 5926912,
    "pageFaults": 4092,
    "majorPageFaults": 0
  }
}
],
"cpu": {
  "time": "2020-08-23T02:29:43Z",
  "usageNanoCores": 11699776,
  "usageCoreNanoSeconds": 654543859283
},
"memory": {
  "time": "2020-08-23T02:29:43Z",
  "availableBytes": 8796087861248,
  "usageBytes": 7589888,
  "workingSetBytes": 5156864,
  "rssBytes": 5914624,
  "pageFaults": 0,
  "majorPageFaults": 0
}
}
]
}

```

Before you deploy metrics-server, you have to make sure that you deploy it on the node which has apiserver deployed on. In my case, that is my master node. As a consequence, I need to make master node

schedulable by the following command:

```
kubectl taint nodes --all node-role.kubernetes.io/master-
```

Then, in the deployment.yaml file, I must specify that metrics-server is deployed on master node.

(I used hostname as the marked label.)

In **metrics-server-deployment.yaml**

```
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/hostname
                operator: In
                values:
                  - charlie-latest
```

Also, you need to check the metrics-server version if it is the latest one. Since the KubeEdge's kubelet is not the original, it uses the different port rather than the original kubelet. That is, metrics-server need to have the ability of searching other port automatically.

Prepare you image first.

Git clone latest metrics server repository:

```
git clone https://github.com/kubernetes-sigs/metrics-server.git
```

Get to the metrics server directory:

```
cd metrics-server
```

Make the docker image:

```
make container-arm64
```

Check if you have this docker image:

```
docker images
```

```
root@raspberrypi:/home/pi# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
SIZE			
gcr.io/k8s-staging-metrics-server/metrics-server-arm64	6d92704c5a68cd29a7a81bce68e6c2230c7a6912	11439f12da24	4 days ago
54.6MB			
<none>	<none>	cafa4702d453	4 days ago
1.41GB			
kubeedge-temperature-mapper	test	c1fd304a4885	12 days ago
389MB			
grafana/grafana	latest	9f31b3fc1ea3	2 weeks ago
149MB			
prom/prometheus	latest	72d6dda76d0b	2 weeks ago
124MB			
gcr.io/k8s-staging-metrics-server/metrics-server-arm	master	3289eecc7ec	2 weeks ago
50MB			
weaveworks/weave-npc	2.7.0	2d47a5fd0000	2 weeks ago
36.7MB			
weaveworks/weave-kube	2.7.0	f58a4b249316	2 weeks ago
90.1MB			
k8s.gcr.io/kube-proxy	v1.18.6	b732e226e0f1	5 weeks ago
97.7MB			

Make sure you change the tag of image by using its IMAGE ID to be compactable with image name in yaml file.

docker tag 11439f12da24 metrics-server-kubeedge:latest

Metrics server works finally!!!!!!!]

```
root@charlie-latest:/var/log/kubeedge# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
charlie-latest	Ready	master	12d	v1.18.6
charlie-test	Ready	<none>	12d	v1.18.6
raspberrypi	Ready	agent,edge	12d	v1.18.6-kubeedge-v1.4.0

```
root@charlie-latest:/var/log/kubeedge# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
temperature-mapper-79f4989f65-5tswz	1/1	Running	0	16h

```
root@charlie-latest:/var/log/kubeedge# kubectl top nodes
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
charlie-latest	459m	7%	4608Mi	59%
charlie-test	126m	3%	1710Mi	45%
raspberrypi	204m	5%	751Mi	9%

```
root@charlie-latest:/var/log/kubeedge# kubectl top pods
```

NAME	CPU(cores)	MEMORY(bytes)
temperature-mapper-79f4989f65-5tswz	12m	4Mi

Hardware: Raspberry Pi 4

KubeEdge version: KubeEdge v1.4.0

Kubernetes version: v1.18.6

Running example: kubeedge-temperature-demo

Metrics-server version: latest-master-branch (Automatically searching port is 0.4.0 feature)

Raspberry Pi 4:

Performance units:

CPU: cores → usage nano cores

Memory: MiB (mebibytes) → 2^{20} Bytes → 1,048,576 Bytes

Q&A in slack channel

1. What I am trying to figure out here is that is there any latest version of metrics-server which is built for "ARM"?
Latest image is available here gcr.io/k8s-staging-metrics-server/metrics-server:master It should be multi arch, so ARM should work with it.
2. Is there any other way to come-across this network issue?
Each network issue is very different, you should check `192.168.0.128:10250` is the correct address where kubelet should be available and debug why it's not accessible
3. Are there any other tools to monitor the metrics without the network configuration?
Depends what metrics do you want, metrics server serves very specific purpose of autoscaling based on CPU and Memory. You can learn more here <https://github.com/kubernetes-sigs/metrics-server/blob/master/README.md>



serathius 4 hours ago

Ad 1 Latest image is available here gcr.io/k8s-staging-metrics-server/metrics-server:master It should be multi arch, so ARM should work with it.

Ad 2 Each network issue is very different, you should check `192.168.0.128:10250` is the correct address where kubelet should be available and debug why it's not accessible

Ad 3 Depends what metrics do you want, metrics server serves very specific purpose of autoscaling based on CPU and Memory. You can learn more here <https://github.com/kubernetes-sigs/metrics-server/blob/master/README.md>

Ref:

Git clone repository, certificates generation file, compilation process:

<https://docs.kubeedge.io/en/v1.2.0/setup/setup.html>

KubeEdge latest document:

<https://docs.kubeedge.io/en/latest/pdf/>

Slack channel for metrics-server in k8s:

https://kubernetes.slack.com/archives/C20HH14P7/p1597942669003600?thread_ts=1597797376.027100&cid=C20HH14P7

How to make k8s master node schedulable:

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

<https://www.thegeekdiary.com/how-to-schedule-master-node-running-pod-service-as-a-worker-node/>

KubeEdge example:

<https://github.com/kubeedge/examples/tree/master/kubeedge-temperature-demo>

Wikipedia-Memory-Units:

<https://zh.wikipedia.org/wiki/Mebibyte>

